Tesora

Frontier AI for Actuaries.

AI helps actuaries perform complex, multi-step actuarial tasks with every step documented and reproducible.

THE WORKBENCH

One collaborative workspace for pricing, reserving, and filing

For all actuaries at admitted and specialty carriers, MGAs, and reinsurers.

August 2026

CONFIDENTIAL

Actuarial work is not one workflow. It is a **web of intricate, connected tasks.**

The more steps an actuary delegates to AI, the less certain the output. Tesora traces and documents every step, so the actuary stays in control.

Philo Bishay · 1st
Cofounder & Head of Product @ Tesora | Actuarial | Wharton MBA
2w ·

Every actuary dreads one question, "What do actuaries do?" I am happy to say we are no closer to an answer, but if you want to avoid further questions, send this:

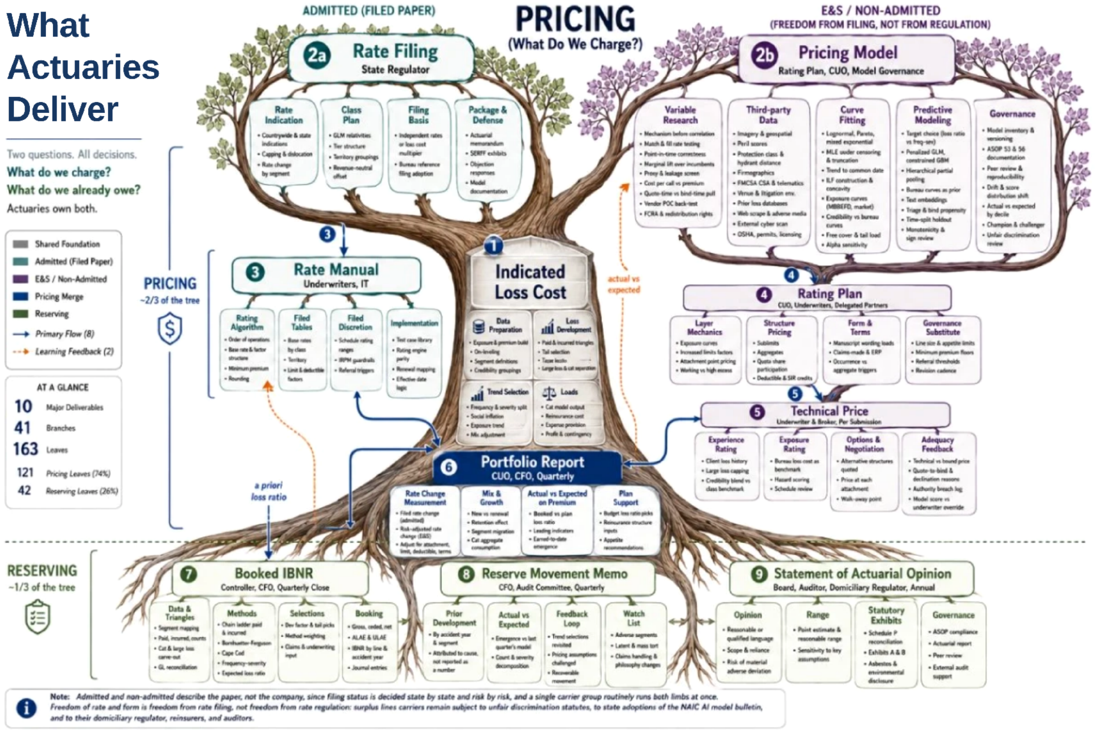

Federico Reyes Gómez and 675 others · 23 comments · 28 reposts

Over 150 workflows are core to the actuarial function

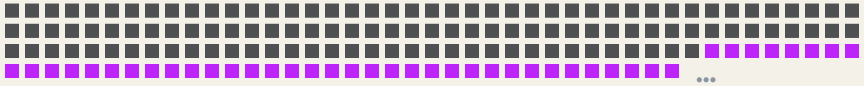

■ 121 pricing ■ 42 reserving — one mark per workflow

Ingest — Triangles — Trend - - ✕ - - Indication — Filing

One weak handoff and everything downstream is unusable.

One shot, and you get a number nobody can trace.

- A rate filing or a booked reserve needs dozens of these correct, in sequence, each one feeding the next.
- Automate a single step and the chain still breaks at the handoff after it.

One agentic harness across all of it

A system of agents keeps every task linked and every value cited, from ingest to filing. That orchestration is the hard part to build, and the harder part to replicate.

No single agent does the whole job, and none of it reaches you unchecked.

We watched how actuaries actually do the work, mapped it into sub-agents, and wrote evals for each one, with audit agents sitting behind them. In a domain this intricate, that is what delivers a business outcome without a hallucinated number or a botched handoff between steps.

A REQUEST ARRIVES

1 **The coordinator breaks the request into tracked work**

It picks the specialists, runs them in parallel, and reports back.

2 **A specialist agent owns each domain end to end**

Each returns a finished deliverable in the format that domain expects.

Pricing and experience
Loss ratios, trend, and frequency-severity, as a documented report.

Reserving
Loss development, IBNR, and triangles, as a formatted exhibit.

Rating models
Builds, validates, and tests a rater from its filing.

Document extraction
Filings, loss runs, and PDFs into clean tables, cell by cell.

Audit your existing work
Reads a spreadsheet, document, or formula you already have and reports what is wrong.

3 **Review agents check the work before you do**

Read access only. **They report findings, and never edit the work.**

Actuarial expert
Method questions, answered from the standards library.

Citation check
Every figure hyperlinks to its page. No link, no value.

Value tracing
Each number ties to a document, or to a sum of ones that do.

Coverage check
Reads what was not cited, to catch anything skipped.

YOUR ACTUARY REVIEWS AND SIGNS OFF

Rate changes in days, not months.

Source a rater from a SERFF filing, an Excel file, or a script. Tesora rebuilds it, tests it, and deploys it whenever underwriting is ready.

BEFORE: BY HAND

4-6 months

From a requested change to a filed, live rater, and where the months actually go

01	**Spec the change** Actuary writes up the revised rates and factors	DAYS
02	**IT ticket, then the queue** The change waits for developer capacity	WEEKS
03	**Re-code the rater** Engineering reimplements the logic by hand	WEEKS
04	**Rebuild exhibits, 50 states** Each state reformatted by hand for SERFF	WEEKS
05	**Manual QA** No test suite, so regressions are caught by eye	WEEKS
06	**Compliance sign-off** Actuarial and legal review the filing	DAYS
07	**File and answer objections** Back and forth with each department of insurance	WEEKS
08	**Release window** Deploy only in the next scheduled window	WEEKS

AFTER: Tesora

1-2 days

Deploy to a GUI, an API, or back into Excel.

You give: a plain-language change and the source rater, from SERFF, Excel, or a script.

You get: a tested rater with a working GUI, sample ratings, a versioned API, or an Excel file your team can open.

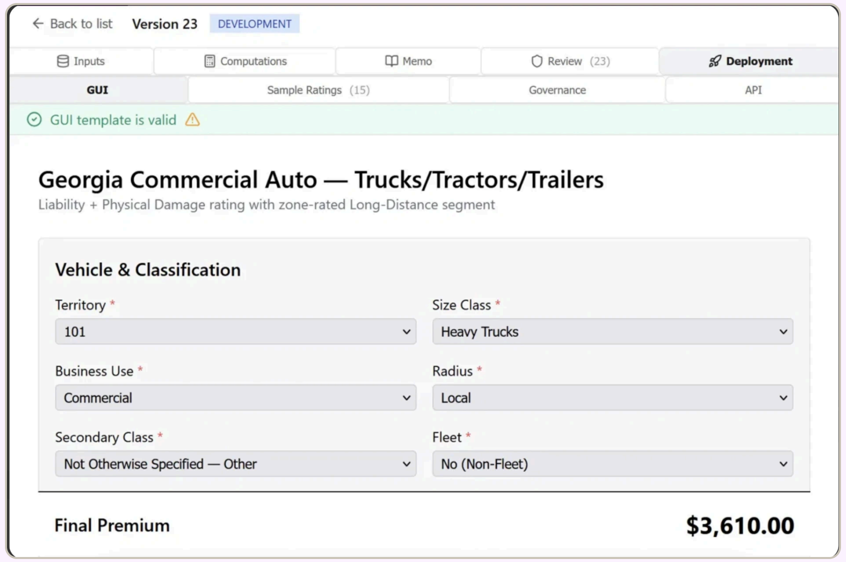

The triangle is squared before you sit down.

A quarterly close moves between actuarial, IT, and finance, and most of the calendar goes to assembling the data rather than judging it.

BEFORE: BY HAND

~2 weeks

From close to a booked IBNR, and who is waiting on whom

01	**Request the data** Actuarial files a ticket for the quarter's extract	IT
02	**Wait on the extract** The pull sits behind other reporting work	DAYS
03	**Reconcile to the ledger** Paid and incurred tied back to finance, by hand	FINANCE
04	**Rebuild the triangles** Segment mapping and large-loss carve-outs redone each quarter	DAYS
05	**Select factors and methods** Chain ladder, Bornhuetter-Ferguson, Cape Cod, tail	ACTUARIAL
06	**Rebuild the exhibits** Roll-forward and movement memo reformatted by hand	DAYS
07	**Peer review** A second actuary retraces the work from scratch	DAYS
08	**Book and explain** Answer the CFO, the auditor, and the board	DAYS

AFTER: Tesora

3-4 hours

Triangles squared on ingest, so the quarter starts at the selection

You give: the loss run, as it comes out of the system.

You get: reconciled triangles, method comparisons, and an IBNR with every value traced to source.

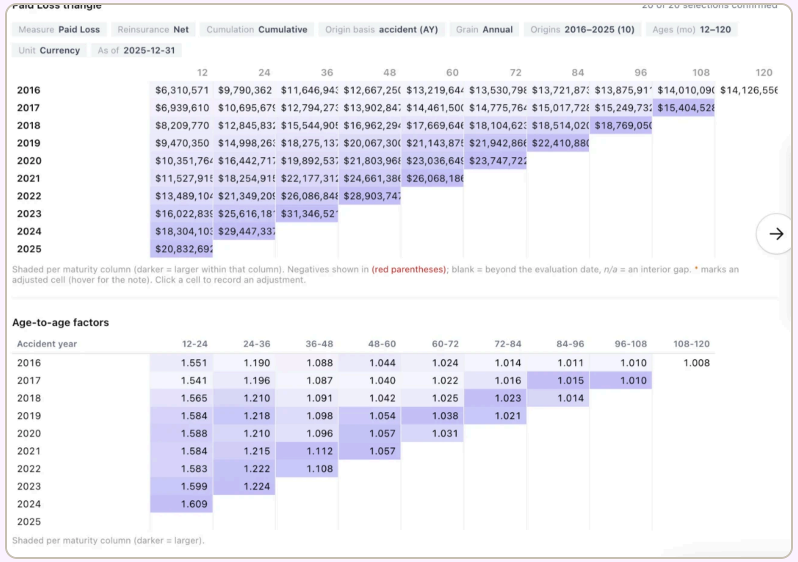

Paid Loss triangle

Measure Paid Loss · Reinsurance Net · Cumulation Cumulative · Origin basis accident (AY) · Grain Annual · Origins 2016–2025 (10) · Ages (mo) 12–120 · Unit Currency · As of 2025-12-31

	12	24	36	48	60	72	84	96	108	120
2016	$6,310,571	$9,790,362	$11,646,94	$12,667,25	$13,219,64	$13,530,79	$13,721,87	$13,875,91	$14,010,09	$14,126,55
2017	$6,939,610	$10,695,67	$12,794,27	$13,902,84	$14,461,50	$14,775,76	$15,017,72	$15,249,73	$15,404,52	
2018	$8,209,770	$12,845,83	$15,544,90	$16,962,29	$17,669,64	$18,104,62	$18,514,02	$18,769,05		
2019	$9,470,350	$14,998,26	$18,275,13	$20,067,30	$21,143,87	$21,942,86	$22,410,88			
2020	$10,351,76	$16,442,71	$19,892,53	$21,803,96	$23,036,64	$23,747,72				
2021	$11,527,91	$18,254,91	$22,177,31	$24,661,38	$26,068,18					
2022	$13,489,10	$21,349,20	$26,086,84	$28,903,74						
2023	$16,022,83	$25,616,18	$31,346,52							
2024	$18,304,10	$29,447,33								
2025	$20,832,69									

Shaded per maturity column (darker = larger within that column). Negatives shown in (red parentheses); blank = beyond the evaluation date, *n/a* = an interior gap. * marks an adjusted cell (hover for the note). Click a cell to record an adjustment.

Age-to-age factors

Accident year	12-24	24-36	36-48	48-60	60-72	72-84	84-96	96-108	108-120
2016	1.551	1.190	1.088	1.044	1.024	1.014	1.011	1.010	1.008
2017	1.541	1.196	1.087	1.040	1.022	1.016	1.015	1.010	
2018	1.565	1.210	1.091	1.042	1.025	1.023	1.014		
2019	1.584	1.218	1.098	1.054	1.038	1.021			
2020	1.588	1.210	1.096	1.057	1.031				
2021	1.584	1.215	1.112	1.057					
2022	1.583	1.222	1.108						
2023	1.599	1.224							
2024	1.609								
2025									

Shaded per maturity column (darker = larger).

Rebuild what a competitor charges, then **price your book through it**.

A filing gives you tables, rules, and inputs rather than a working model, and the order of operations is rarely spelled out. Turning that into something you can run your own risks through is the work.

BEFORE: REBUILT BY HAND FROM SERFF

~1 week

From a competitor filing to a usable read, and where it goes wrong

01	**Pull the filings**	An analyst searches SERFF, state by state	HOURS
02	**Read the rate manual**	Rules, factors, and order of operations, by eye	DAYS
03	**Rekey the rate tables**	Hundreds of tables retyped into a spreadsheet	DAYS
04	**Rebuild the algorithm**	The logic is inferred, then argued over	DAYS
05	**Price a sample**	A dozen test risks, never the real book	HOURS
06	**Compare by hand**	Deltas built in a one-off workbook nobody owns	DAYS
07	**Explain the gaps**	Each difference traced back to a single factor	DAYS
08	**The read goes stale**	They filed again while you were rebuilding	ONGOING

AFTER: Tesora

10-15 minutes

An executable competitor rater, and your whole book priced through it

You give: a competitor, a state, and a line. Or the SERFF filing itself.

You get: their rating logic running against your risks, every factor cited to the page it came from.

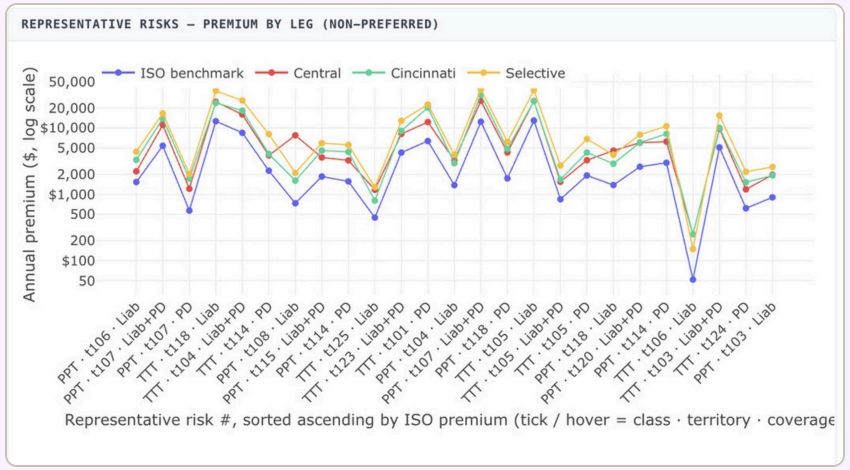

Built from public rate filings, and any additional data sources you provide.

The **full context** is baked into the work.

Every session holds the data it used, the assumptions behind it, and the reasoning that got there. Reopen it years later and nothing has to be reconstructed, because the context was never separate from the answer.

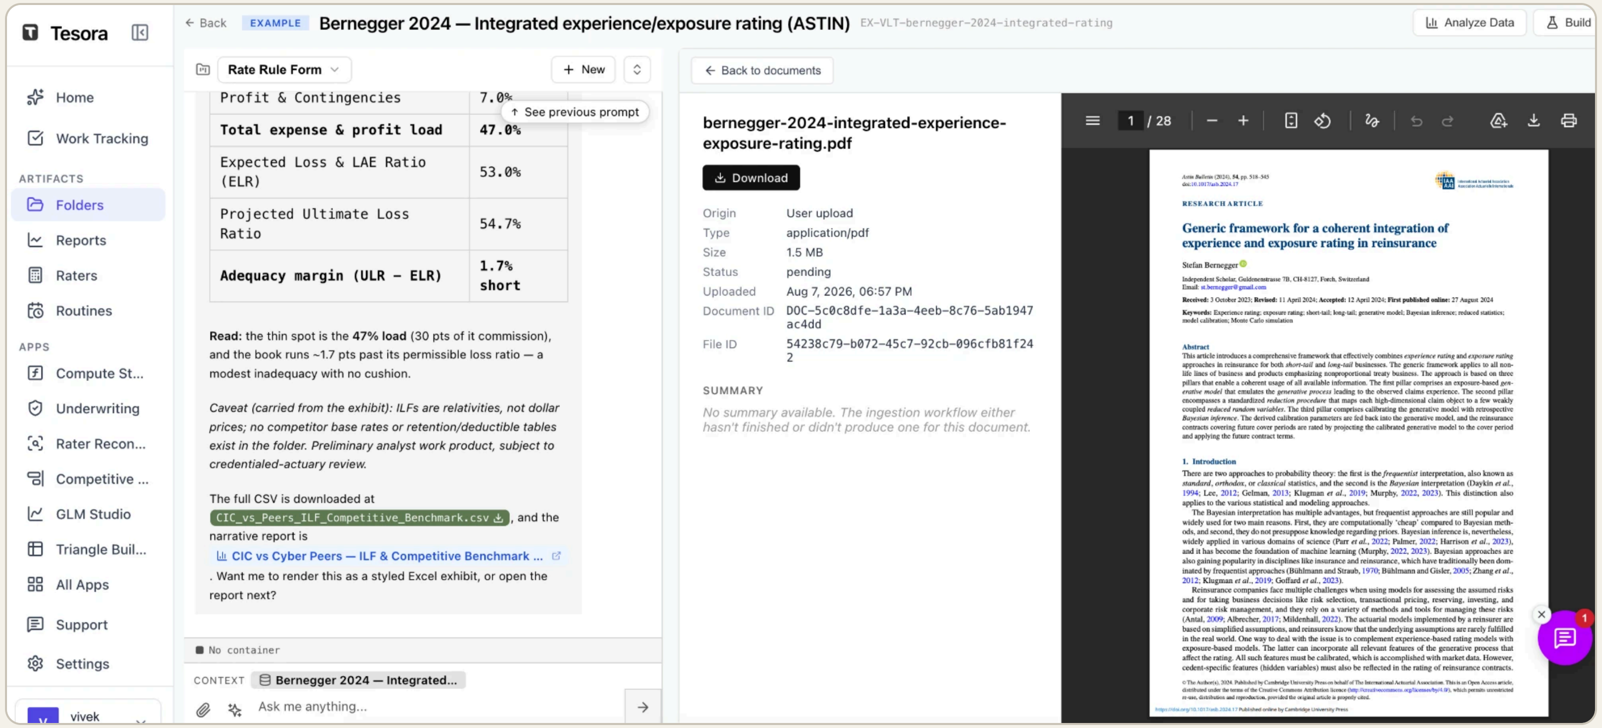

Every request from the platform pings our chief actuary and CTO. Support is 24/7, agents and people together, so you are never left with a bot. A real person answers within two to three minutes, and we have had no unplanned downtime to date, which is **model routing** doing its job.

Implementation

The models are wired into your loss runs, policy systems, and ledger, running inside your own environment.

Governance

Author, reviewer, and approver are separate identities, and the trail survives every promotion. Finance and IT find what they need without waiting on an actuary.

Compliance

Built to ASOP from the start, and SOX-aligned for carriers with public-company parents. Same input, same answer, every time.

Continuity

When an actuary leaves, their assumptions, sources, and reasoning stay in the session. The next person picks up the thread.

What's the best way to implement AI?

Your implementation of AI today will dictate its value, translatability, and reproducibility going forward.

	Actuarial consultants Milliman, WTW, Oliver Wyman	Your team with AI In-house, on ChatGPT or Claude	Actuarial software hyperexponential, Earnix, Akur8	Tesora Frontier models, in your format
Turnaround	Weeks to months, in their queue	Quick drafts, still checked by hand	Fast and repeatable once it is live	Minutes to days, in your workflow
Setup & fit	A fresh engagement scoped each time	No one knows your book better	Heavy implementation on their platform	60-day build in your VPC, on your own data
Your raters & filings	Generic deliverables in their template	Your format, rebuilt by hand each time	Rebuilt inside their system, their way	Your raters, versioned and ready to deploy
Source & audit	Independent, credible workpapers	No source to trace, and the answer drifts	Logged inside the platform	Cell-level citations, append-only log
Scope	Whatever you scope and pay for	One step at a time, gaps in between	Pricing or reserving, rarely filing too	Pricing, reserving, filing, and close in one place

● Full capability ○ Gap

A team built at the intersection of frontier AI research and decades of actuarial experience, the two disciplines this problem demands.

Vivek Rao

CEO

Former private-equity investor and McKinsey consultant focused on P&C insurance.

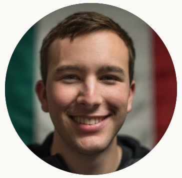

Federico Reyes

CTO

Stanford AI researcher. Graph neural nets and relational data at Kumo AI, acquired by Nvidia.

Philo Bishay

CPO & CHIEF ACTUARY

Former chief actuary at W. R. Berkley E&S. Capital modeler at **Gen Re**; built and sold a brokerage.

Michael Palmich

HEAD OF SALES

25+ years in the actuarial software space, most recently at Akur8 and Milliman.

Eric Lam

FOUNDING ENGINEER

MIT for CS, physics, and math, with **three actuarial exams** done.

ADVISORY BOARD

Ethan Genteman

FCAS, MAAA

Chief Actuary, **Palomar** (NYSE: PLMR)

Mario DiCaro

FCAS, CERA

Head of Capital Modeling, **Tokio Marine HCC**

Joanne Chen

General Partner, **Foundation Capital**

BACKED BY

foundation capital — Leading AI investors, backing enterprise founders since 1995

Y Combinator S25 — The accelerator behind Stripe and Airbnb

Proud sponsors of CAS

Tesora

Above all, we want to help you win.

We are a thought partner first. This technology is powerful, and it only helps once it is opinionated: built with a real point of view on how the work gets done, how it gets checked, and how it gets defended when a regulator asks.

We are here to push the frontier forward with the right partners, not to sell you a license. A law firm does the legal work rather than selling software to lawyers, and we hold ourselves to that standard: **we take ownership of the outcome, not just the tooling**. Every deployment starts with a proof of concept anchored on your own success metrics, and we work with your team to find where the workbench sets the shared standard first.

Let's set up a call

sales@tesora.ai

DE-RISKED BY DESIGN

- ✓ It runs in your own environment, in a dedicated tenant. No cross-tenant inference.
- ✓ Your rate plans, filings, and books are never used to train shared models.
- ✓ AES-256 at rest, TLS 1.3 in transit, and your own SSO.
- ✓ We clear your security review before any work touches a filing.

SOC 2 Type II. Independent auditors review us annually, and our trust center shows the continuous testing live. SOX-aligned against both actuarial and IT standards.

Appendix

Per actuary, the recurring month goes from **150 hours to about 60**.

Most of the month disappears into recurring work: the data pulls, the reconciliations, the model updates.
Tesora compresses that mechanical layer, so the hours come back as judgment rather than assembly.

Where one actuary's recurring month goes (reclaimed / still hands-on; scale 10, 20, 30, 40 hrs)

Workflow	Hours	Gain
Data extraction & reconciliation	39 → 4	10×
Pricing model updates	28 → 7	4×
Loss trend analysis	23 → 12	2×
Reserve & loss development	21 → 12	1.7×
Management & board reporting	17 → 10	1.7×
Reinsurance program analysis	12 → 4	3×
Peer review & documentation	10 → 8	1.3×
Monthly total, per actuary (93 hrs reclaimed, 57 hands-on)	150 → ~60	2.5×

~90 hours back per actuary, per month

A 2.5× blended throughput gain across the recurring month, with the same team and the same sign-off.

Smarter, not harder

Time freed for strategy and selections rather than assembly.

Unlock your actuaries

More products, more territories, and more actuarial capability from the same team.

Hours are per actuary and reflect a typical ~150 hour recurring workload within a roughly 160 hour month, with per-workflow gains from Tesora's engagement framework. Figures are modeled and confirmed against measured results during a build.

We hand the mechanical hours back to your team as actuarial judgment, freeing them to focus where it counts. Capacity stops being the ceiling.

We run on top of Claude, and we are model agnostic.

The model is the engine, not the product. What an actuary needs around it is infrastructure: the data wired in, the workflow decided, the cost controlled, and the record kept. That is the part we build, and we build it for you.

We guide the model	A raw model answers anything, confidently. Ours is held to an actuarial method: fit frequency and severity separately, test against your own history, and stop to ask before selecting.
We wire the infrastructure	Your loss runs, policy systems, filings, and ledger are connected once and stay connected, so every study reads the same reconciled book. A chat window starts empty every time.
The workflow is opinionated	Reserving and pricing have a right order of operations, and we encode it. Chain ladder, Bornhuetter-Ferguson, and Cape Cod run side by side, so the selection stays your actuary's call.
We manage the cost	Running a real book through a frontier model naively is slow and expensive. Extraction runs on a small model, and the judgment calls go to the strongest one.

We exist to build your actuarial infrastructure, not to serve the models. When a better one ships you get it, and the system around your book is what stays.

We use the best models. You keep the value.

Every firm is working out its AI spend internally right now. Our team of AI researchers and former investors knows what these systems cost, and how to optimize model spend per business outcome.

Right model, right step	Each step goes to the model that gets it right for the least spend. Extraction runs small, judgment runs strong, and none of that routing is your problem to manage.
Our incentive is the outcome, not the token count	A model vendor is never going to tell you to buy less of what it sells. Our platform has to deliver the result, so keeping model spend down is work we have to solve for ourselves rather than work we are merely free to do. **You are never left at the mercy of a rogue agent burning through tokens.**
Latest models, always up	Your workflows, data wiring, and audit trail live with us, so the strongest model arrives without a migration or a new vendor. A system built in-house on one provider stops when that provider does. We route across them, so **servicing an outage is our problem rather than your stopped quarter**.
Context mapping	We map your data, formats, and quirks against the work they feed, so we know intimately which model fits where. That map deepens the longer we work with your team, and it carries across every upgrade.
Audits matter more, not less	Better models take on more complex work, which raises rather than lowers the bar on proving it. Every value stays cited and reproducible, and **the advantage compounds for whoever moves first**.