



Frontier AI for Actuaries.

AI helps actuaries perform complex, multi-step actuarial tasks with every step documented and reproducible.

THE WORKBENCH

One collaborative workspace for pricing, reserving, and filing


For all actuaries at admitted and specialty carriers, MGAs, and reinsurers.

August 2026
CONFIDENTIAL



Actuarial work is not one workflow. It is a **web of intricate, connected tasks**.

The more steps an actuary delegates to AI, the less certain they can be of the output. Tesora enables full tracing, documentation, and audit, so the actuary remains in control.



ONE ACTUARIAL FUNCTION, MAPPED

10 major deliverables

41 branches

163 leaves, each a real decision

~2/3
pricing

~1/3
reserving

Every leaf is a judgment someone signs for. Automate one and the other 162, with the handoffs between them, are still yours.



One agentic harness across all of it

A system of agents keeps every task linked and every value cited, from ingest to filing. That orchestration is the hard part to build, and the harder part to replicate.



What's the best way to **implement AI**?

Your implementation of AI today will dictate its value, translatability, and reproducibility going forward.

	Actuarial consultants Milliman, WTW, Oliver Wyman	Your team with AI In-house, on ChatGPT or Claude	Actuarial software hyperexponential, Earnix, Akur8	T Tesora Frontier models, in your format
Turnaround	✗ Weeks to months, in their queue	≈ Quick drafts, still checked by hand	✓ Fast and repeatable once it is live	✓ Minutes to days, in your workflow
Setup & fit	≈ A fresh engagement scoped each time	✓ No one knows your book better	✗ Heavy implementation on their platform	✓ 60-day build in your VPC, on your own data
Your raters & filings	✗ Generic deliverables in their template	≈ Your format, rebuilt by hand each time	≈ Rebuilt inside their system, their way	✓ Your raters, versioned and ready to deploy
Source & audit	✓ Independent, credible workpapers	✗ No source to trace, and the answer drifts	≈ Logged inside the platform	✓ Cell-level citations, append-only log
Scope	≈ Whatever you scope and pay for	≈ One step at a time, gaps in between	≈ Pricing or reserving, rarely filing too	✓ Pricing, reserving, filing, and close in one place



Strength



Partial or conditional



Gap



Rate changes in **days**, not months.

Source a rater from a SERFF filing, an Excel file, or a script. Tesora rebuilds it, tests it, and deploys it whenever underwriting is ready.

BEFORE · BY HAND

4-6 months

From a requested change to a filed, live rater, and where the months actually go

- | | | |
|----|---|-------|
| 01 | Spec the change
Actuary writes up the revised rates and factors | DAYS |
| 02 | IT ticket, then the queue
The change waits for developer capacity | WEEKS |
| 03 | Re-code the rater
Engineering reimplements the logic by hand | WEEKS |
| 04 | Rebuild exhibits, 50 states
Each state reformatted by hand for SERFF | WEEKS |
| 05 | Manual QA
No test suite, so regressions are caught by eye | WEEKS |
| 06 | Compliance sign-off
Actuarial and legal review the filing | DAYS |
| 07 | File and answer objections
Back and forth with each department of insurance | WEEKS |
| 08 | Release window
Deploy only in the next scheduled window | WEEKS |

AFTER · ON TESORA

1-2 days

Deploy as a web interface or API.

- | | |
|----------|---|
| YOU GIVE | A plain-language change and the source rater, from SERFF, Excel, or a script. |
| YOU GET | A tested rater with a working GUI, sample ratings, and a versioned API. |



The triangle is squared before you sit down.

A quarterly close moves between actuarial, IT, and finance, and most of the calendar goes to assembling the data rather than judging it.

BEFORE · BY HAND

~2 weeks

From close to a booked IBNR, and who is waiting on whom

01 Request the data

Actuarial files a ticket for the quarter's extract

IT

02 Wait on the extract

The pull sits behind other reporting work

DAYS

03 Reconcile to the ledger

Paid and incurred tied back to finance, by hand

FINANCE

04 Rebuild the triangles

Segment mapping and large-loss carve-outs redone each quarter

DAYS

05 Select factors and methods

Chain ladder, Bornhuetter-Ferguson, Cape Cod, tail

ACTUARIAL

06 Rebuild the exhibits

Roll-forward and movement memo reformatted by hand

DAYS

07 Peer review

A second actuary retraces the work from scratch

DAYS

08 Book and explain

Answer the CFO, the auditor, and the board

DAYS


AFTER · ON TESORA

3-4 hours

Triangles squared on ingest, so the quarter starts at the selection

YOU GIVE The loss run, as it comes out of the system.

YOU GET Reconciled triangles, method comparisons, and an IBNR with every value traced to source.





Work you can reopen years later, by whoever is here then.

AI sessions save your work so you can revisit, reproduce, or audit it many years later. The reasoning stays with the company rather than with the person who did it.

Tesora | Back | EXAMPLE | Bernegger 2024 — Integrated experience/exposure rating (ASTIN) | EX-VLT-bernegger-2024-integrated-rating | Analyze Data | Build

Rate Rule Form | + New

Profit & Contingencies	7.0%	↑ See previous prompt
Total expense & profit load	47.0%	
Expected Loss & LAE Ratio (ELR)	53.0%	
Projected Ultimate Loss Ratio	54.7%	
Adequacy margin (ULR - ELR)	1.7%	short

Read: the thin spot is the **47% load** (30 pts of it commission), and the book runs ~1.7 pts past its permissible loss ratio — a modest inadequacy with no cushion.

Caveat (carried from the exhibit): ILFs are relativities, not dollar prices; no competitor base rates or retention/deductible tables exist in the folder. Preliminary analyst work product, subject to credentialed-actuary review.

The full CSV is downloaded at `CIC_vs_Peers_ILF_Competitive_Benchmark.csv`, and the narrative report is [CIC vs Cyber Peers — ILF & Competitive Benchmark ...](#). Want me to render this as a styled Excel exhibit, or open the report next?

CONTEXT | Bernegger 2024 — Integrated... | Ask me anything...

bernegger-2024-integrated-experience-exposure-rating.pdf | Download

Origin: User upload
Type: application/pdf
Size: 1.5 MB
Status: pending
Uploaded: Aug 7, 2026, 06:57 PM
Document ID: DOC-5c0c8dfe-1a3a-4eeb-8c76-5ab1947ac4dd
File ID: 54238c79-b072-45c7-92cb-096c7b81f242

SUMMARY
No summary available. The ingestion workflow either hasn't finished or didn't produce one for this document.

1. Introduction
There are two approaches to probability theory: the first is the *frequentist* interpretation, also known as *standard, orthodox, or classical* statistics, and the second is the *Bayesian* interpretation (DePistis et al., 1994; Lee, 2012; Gelman, 2013; Klugman et al., 2019; Murphy, 2022, 2023). This distinction also applies to the various statistical and modeling approaches.

The Bayesian interpretation has multiple advantages, but frequentist approaches are still popular and widely used for two main reasons. First, they are computationally 'cheap' compared to Bayesian methods, and second, they do not presuppose knowledge regarding priors. Bayesian inference is, nevertheless, widely applied in various domains of science (Parr et al., 2022; Palmer, 2022; Harrison et al., 2023), and it has become the foundation of machine learning (Murphy, 2022, 2023). Bayesian approaches are also gaining popularity in disciplines like insurance and reinsurance, which have traditionally been dominated by frequentist approaches (Bühlmann and Straub, 1970; Bühlmann and Gisler, 2005; Zhang et al., 2012; Klugman et al., 2019; Goffard et al., 2023).

Reinsurance companies face multiple challenges when using models for assessing the assumed risks and for taking business decisions like risk selection, transactional pricing, reserving, investing, and corporate risk management, and they rely on a variety of methods and tools for managing these risks (Antal, 2009; Albrecher, 2017; Mildenhall, 2022). The actuarial models implemented by a reinsurer are based on simplified assumptions, and reinsurers know that the underlying assumptions are rarely fulfilled in the real world. One way to deal with the issue is to complement experience-based rating models with exposure-based models. The latter can incorporate all relevant features of the generative process that affect the rating. All such features must be calibrated, which is accomplished with market data. However, cedent-specific features (hidden variables) must also be reflected in the rating of reinsurance contracts.

© The Author(s), 2024. Published by Cambridge University Press on behalf of The International Actuarial Association. This is an Open Access article, distributed under the terms of the Creative Commons Attribution licence (http://creativecommons.org/licenses/by/4.0/), which permits unrestricted re-use, distribution and reproduction, provided the original article is properly cited.

A When an actuary leaves

Their assumptions, sources, and reasoning stay in the session. The next person picks up the thread instead of rebuilding it.

F Finance

Ties a booked number back to the ledger without waiting on an actuary to explain it.

IT IT and data

Sees which files fed which number, so a schema change has a known blast radius.

R Auditors and regulators

Follow any value to its source document, years after the filing closed.

One shared record of how the work was done, and it outlasts the people who did it.



We run on top of Claude, and we are model agnostic.

The model is the engine, not the product. What an actuary needs around it is infrastructure: the data wired in, the workflow decided, the cost controlled, and the record kept. That is the part we build, and we build it for you.

01 We guide the model

A raw model will answer anything, confidently. Ours is held to an actuarial method: which factors to consider, which checks to run, when to stop and ask. That guidance is the difference between a plausible answer and a defensible one.

02 We wire the infrastructure

Your loss runs, policy systems, filings, and ledger are connected once and stay connected. A chat window starts empty every time; the workbench already knows your book.

03 The workflow is opinionated

Reserving and pricing have a right order of operations. We encode it, so the work arrives in the shape an actuary expects rather than whatever the prompt happened to produce.

04 We manage the token cost

Running a real book through a frontier model naively is slow and expensive. We route each step to the cheapest model that is still correct, so the economics work at production scale.



We exist to build **your actuarial infrastructure**, not to serve the models or ourselves. When a better model ships, you get it. What stays is the system we built around your book.



It makes us **even better.**

Better models accelerate the product underneath you, compounding your advantage with every release.



Every release lifts you

The workbench is model-agnostic and runs on the strongest model available. Each new generation handles your schemas and messy edge cases more easily, so the work gets sharper **at no cost to you.**



We keep learning your systems

The longer we work with your team, the more of your data, formats, and quirks the workbench encodes, and that knowledge carries across every model upgrade.



The cheapest correct path

Every org is running a dozen AI pilots. We stay anchored on one thing: the fastest, most correct, and lowest-token way to run a complex workflow within your audit standards.



Audits matter more, not less

As the model does more of the work, proving it matters more, so no step gets skipped and nothing is taken for granted. Every value stays cited, reproducible, and signed off.



The gap widens

Better models unlock even more complex work, so the advantage compounds. A rate change that took months now takes days, so you can reprice against the market while slower competitors wait on a filing cycle. The firms that move first pull ahead, and **standing still gets more expensive.**

Above all, we want to help you win.

We are a thought partner first. This technology is powerful, and it only helps once it is opinionated: built with a real point of view on how the work gets done and defended.

We are here to push the frontier forward with the right partners, not to sell you a license. Doing right by you comes first, so every deployment starts with a proof of concept anchored on your organization's own success metrics. Some teams are stronger at reserving, others at pricing, and we work with yours to find where the workbench sets the best practices and shared learnings first.

DE-RISKED BY DESIGN

- ✓ It runs in your own environment. We never hold your data.
- ✓ SOC 2 Type II, with the highest data-security controls.
- ✓ We clear your security review before any work touches a filing.

Start a POC discussion

sales@tesora.ai

APPENDIX

The supporting detail.

A What the recurring month looks like, and the ROI math behind it

B The team, and who backs us



The recurring month goes from 150 hours to about 60.

Tesora compresses the mechanical grind, the data pulls, the reconciliations, the model updates, so the same team spends its hours on judgment instead of assembly.



2.5×

the recurring workload, same team and sign-off

The blended throughput gain across the whole recurring month.

Smarter not harder

Time freed for strategy and selections, not assembly

Unlock your actuaries

More products, territories, and actuarial capability

Hours reflect a typical ~150 hour recurring actuarial month, with per-workflow gains from Tesora's engagement framework. Figures are modeled and confirmed against measured results during a build.




A team built at the intersection of frontier AI research and decades of actuarial experience, **the two disciplines this problem demands.**



Vivek Rao

CEO

Former private-equity investor and McKinsey consultant focused on P&C insurance.



Federico Reyes

CTO

Stanford AI researcher. Graph neural nets and relational data at Kumo AI (acquired by Nvidia), and on Google's Document OCR team.



Philo Bishay

CPO & CHIEF ACTUARY

Chief actuary at W. R. Berkley E&S, owning pricing and reserving. Capital modeler at **Gen Re**; built and sold a brokerage.

BACKED BY



Leading AI and enterprise investors since 1995



Y Combinator **S25**

The accelerator behind Stripe and Airbnb

☆ **Proud sponsor of the CAS**