

Frontier AI for Actuaries.

AI helps actuaries perform complex, multi-step actuarial tasks with every step documented and reproducible.

THE WORKBENCH


One collaborative workspace for pricing, reserving, and filing

For all actuaries at admitted and specialty carriers, MGAs, and reinsurers.



Actuarial work is not one workflow. It is a **web of intricate, connected tasks**.

The more steps an actuary delegates to AI, the less certain they can be of the output. Tesora enables full tracing, documentation, and audit, so the actuary remains in control.



One agentic harness across all of it

A system of agents keeps every task linked and every value cited, from ingest to filing. That orchestration is the hard part to build, and the harder part to replicate.

AUDITED END TO END



The same question, and only one answer you can put your name on.

A frontier model can do the math. It cannot show you where the numbers came from, and it will not give you the same answer twice.

A FRONTIER MODEL, ON ITS OWN One answer, no trail

 Claude

Select a loss trend for my GA commercial auto book and give me the indicated rate change.

 Based on typical industry patterns for commercial auto, a loss trend near **6%** is reasonable, which supports a **mid-single-digit** rate change.

! WHY AN ACTUARY CANNOT SIGN THIS

- ! It never saw your data. Nothing here traces to your book.
- ! No record of what it did, so no one can review the steps.
- ! Ask again next quarter and the answer moves.

THE SAME MODELS, INSIDE TESORA Every step on the record

 Loss trend selection

session #4471

- ✓ Pulled your paid and incurred history
↳ GA_CA_losses_2019-2025.xlsx
- ✓ Fit frequency and severity trends separately
↳ exhibit 4, rows 12-31
- ✓ Cross-checked against filed industry trend
↳ SERFF GA-2025-0118
- ✓ Your actuary selected and signed off
↳ reviewed by P. Bishay
- ✓ Rerun it in two years and it reproduces, cell for cell.



What's the best way to **implement AI**?

Your implementation of AI today will dictate its value, translatability, and reproducibility going forward.

	Actuarial consultants Milliman, WTW, Oliver Wyman	Your team with AI In-house, on ChatGPT or Claude	Actuarial software hyperexponential, Earnix, Akur8	T Tesora Frontier models, in your format
Turnaround	✗ Weeks to months, in their queue	≈ Quick drafts, still checked by hand	✓ Fast and repeatable once it is live	✓ Minutes to days, in your workflow
Setup & fit	≈ A fresh engagement scoped each time	✓ No one knows your book better	✗ Heavy implementation on their platform	✓ 60-day build in your VPC, on your own data
Your raters & filings	✗ Generic deliverables in their template	≈ Your format, rebuilt by hand each time	≈ Rebuilt inside their system, their way	✓ Your raters, versioned and ready to deploy
Source & audit	✓ Independent, credible workpapers	✗ No source to trace, and the answer drifts	≈ Logged inside the platform	✓ Cell-level citations, append-only log
Scope	≈ Whatever you scope and pay for	≈ One step at a time, gaps in between	≈ Pricing or reserving, rarely filing too	✓ Pricing, reserving, filing, and close in one place



Strength



Partial or conditional



Gap



Rate changes in **days**, not months.

Source a rater from a SERFF filing, an Excel file, or a script. Tesora rebuilds it, tests it, and deploys it whenever underwriting is ready.

BEFORE · BY HAND

4-6 months

From a requested change to a filed, live rater, and where the months actually go

- | | | |
|----|---|-------|
| 01 | Spec the change
Actuary writes up the revised rates and factors | DAYS |
| 02 | IT ticket, then the queue
The change waits for developer capacity | WEEKS |
| 03 | Re-code the rater
Engineering reimplements the logic by hand | WEEKS |
| 04 | Rebuild exhibits, 50 states
Each state reformatted by hand for SERFF | WEEKS |
| 05 | Manual QA
No test suite, so regressions are caught by eye | WEEKS |
| 06 | Compliance sign-off
Actuarial and legal review the filing | DAYS |
| 07 | File and answer objections
Back and forth with each department of insurance | WEEKS |
| 08 | Release window
Deploy only in the next scheduled window | WEEKS |

AFTER · ON TESORA

1-2 days

Deploy as a web interface or API.

- | | |
|----------|---|
| YOU GIVE | A plain-language change and the source rater, from SERFF, Excel, or a script. |
| YOU GET | A tested rater with a working GUI, sample ratings, and a versioned API. |



The triangle is squared before you sit down.

Reserving usually burns days pulling and reconciling triangles. Tesora does that on ingest, so your time goes to selecting factors.



Reserving study screenshot

Selected IBNR by accident year, chain-ladder ultimates reported against paid, every value traced to the source triangle.



Work you can **reopen years later.**

AI sessions save your work so you can revisit, reproduce, or audit the work many years later.




Document citation screenshot

Tesora reading a source document and citing every value it used, with any caveat surfaced for review.



The recurring month goes from 150 hours to about 60.

Tesora compresses the mechanical grind, the data pulls, the reconciliations, the model updates, so the same team spends its hours on judgment instead of assembly.



2.5×

the recurring workload, same team and sign-off

The blended throughput gain across the whole recurring month.

Smarter not harder

Time freed for strategy and selections, not assembly

Unlock your actuaries

More products, territories, and actuarial capability

Hours reflect a typical ~150 hour recurring actuarial month, with per-workflow gains from Tesora's engagement framework. Figures are modeled and confirmed against measured results during a build.

We hand the mechanical hours back to your team as actuarial judgment, freeing them to focus where it counts. Capacity stops being the ceiling.



A team built at the intersection of frontier AI research and decades of actuarial experience, **the two disciplines this problem demands.**



Vivek Rao

CEO

Former private-equity investor and McKinsey consultant focused on P&C insurance.



Federico Reyes

CTO

Stanford AI researcher. Graph neural nets and relational data at Kumo AI (acquired by Nvidia), and on Google's Document OCR team.



Philo Bishay

CPO & CHIEF ACTUARY

Chief actuary at W. R. Berkley E&S, owning pricing and reserving. Capital modeler at **Gen Re**; built and sold a brokerage.

BACKED BY



Foundation Capital

Leading AI and enterprise investors since 1995



Y Combinator

S25

The accelerator behind Stripe and Airbnb

One team, focused on solving this for every actuary

Start with one workflow. Keep the sign-off.

Bringing AI into actuarial work is a decision you will have to explain, to your management, your auditors, and eventually a regulator. That deserves care, so we do not ask for a leap of faith.

We start narrow: one workflow, one metric your team already tracks, running in your environment on your data. Your actuaries keep the judgment and the sign-off. Every step is documented, so you can check the work rather than trust it.

If it earns a place in your process, we go further together. **If it does not, you have spent a few weeks and gained a clear answer.**

DE-RISKED BY DESIGN

- ✓ It runs in your own environment. We never hold your data.
- ✓ SOC 2 Type II, with the highest data-security controls.
- ✓ We clear your security review before any work touches a filing.
- ✓ Every number stays traceable, so your review is a check, not a leap.

Start a POC discussion

sales@tesora.ai