Tesora

Frontier AI for Actuaries.

AI helps actuaries perform complex, multi-step actuarial tasks with every step documented and reproducible.

THE WORKBENCH

One collaborative workspace for pricing, reserving, and filing

For all actuaries at admitted and specialty carriers, MGAs, and reinsurers.

August 2026

Actuarial work is not one workflow. It is a web of intricate, connected tasks.

The more steps an actuary delegates to AI, the less certain they can be of the output. Tesora enables full tracing, documentation, and audit, so the actuary remains in control.

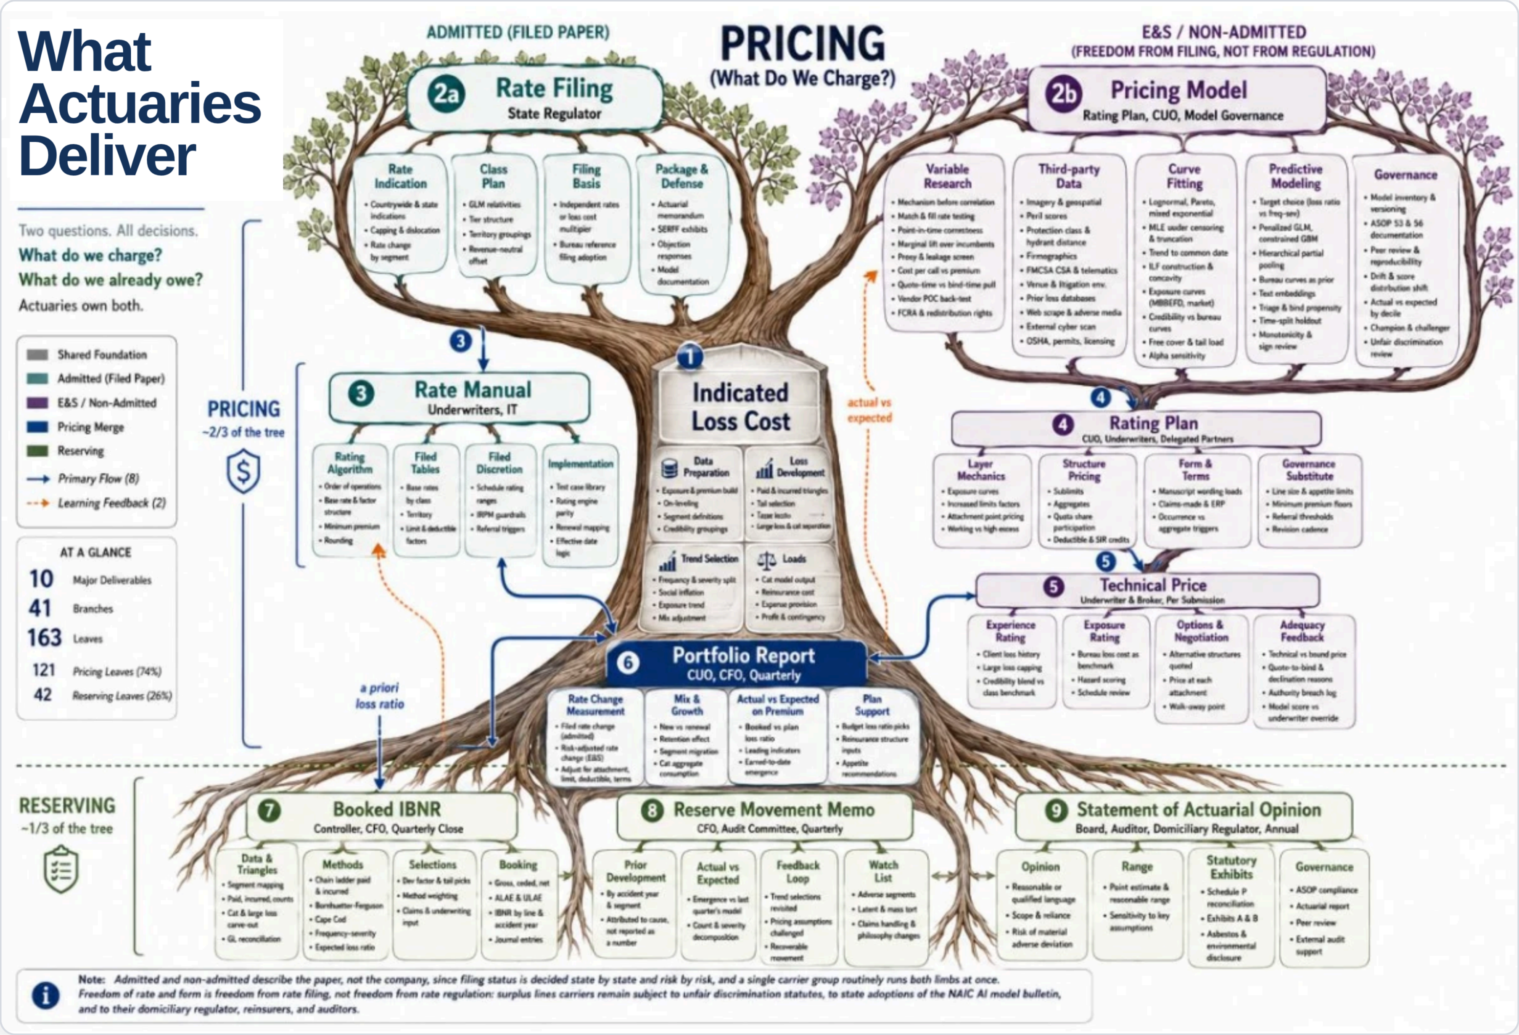

ONE ACTUARIAL FUNCTION, MAPPED

10 major deliverables

41 branches

163 leaves, each a real decision

~2/3	~1/3
pricing	reserving

Every leaf is a judgment someone signs for. Automate one and the other 162, with the handoffs between them, are still yours.

One agentic harness across all of it

A system of agents keeps every task linked and every value cited, from ingest to filing. That orchestration is the hard part to build, and the harder part to replicate.

What's the best way to implement AI?

Your implementation of AI today will dictate its value, translatability, and reproducibility going forward.

	Actuarial consultants Milliman, WTW, Oliver Wyman	Your team with AI In-house, on ChatGPT or Claude	Actuarial software hyperexponential, Earnix, Akur8	Tesora Frontier models, in your format
Turnaround	✗ Weeks to months, in their queue	≈ Quick drafts, still checked by hand	✓ Fast and repeatable once it is live	✓ Minutes to days, in your workflow
Setup & fit	≈ A fresh engagement scoped each time	✓ No one knows your book better	✗ Heavy implementation on their platform	✓ 60-day build in your VPC, on your own data
Your raters & filings	✗ Generic deliverables in their template	≈ Your format, rebuilt by hand each time	≈ Rebuilt inside their system, their way	✓ Your raters, versioned and ready to deploy
Source & audit	✓ Independent, credible workpapers	✗ No source to trace, and the answer drifts	≈ Logged inside the platform	✓ Cell-level citations, append-only log
Scope	≈ Whatever you scope and pay for	≈ One step at a time, gaps in between	≈ Pricing or reserving, rarely filing too	✓ Pricing, reserving, filing, and close in one place

✓ Strength ≈ Partial or conditional ✗ Gap

Rate changes in days, not months.

Source a rater from a SERFF filing, an Excel file, or a script. Tesora rebuilds it, tests it, and deploys it whenever underwriting is ready.

BEFORE · BY HAND

4-6 months

From a requested change to a filed, live rater, and where the months actually go

01	**Spec the change** Actuary writes up the revised rates and factors	DAYS
02	**IT ticket, then the queue** The change waits for developer capacity	WEEKS
03	**Re-code the rater** Engineering reimplements the logic by hand	WEEKS
04	**Rebuild exhibits, 50 states** Each state reformatted by hand for SERFF	WEEKS
05	**Manual QA** No test suite, so regressions are caught by eye	WEEKS
06	**Compliance sign-off** Actuarial and legal review the filing	DAYS
07	**File and answer objections** Back and forth with each department of insurance	WEEKS
08	**Release window** Deploy only in the next scheduled window	WEEKS

AFTER · ON TESORA

1-2 days

Deploy as a web interface or API.

YOU GIVE A plain-language change and the source rater, from SERFF, Excel, or a script.

YOU GET A tested rater with a working GUI, sample ratings, and a versioned API.

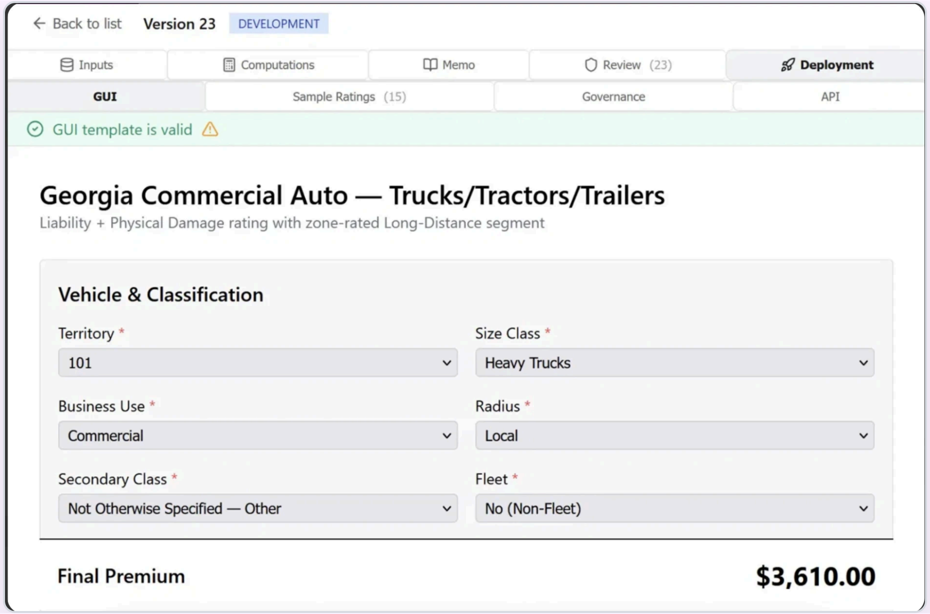

The triangle is squared before you sit down.

A quarterly close moves between actuarial, IT, and finance, and most of the calendar goes to assembling the data rather than judging it.

BEFORE · BY HAND

~2 weeks

From close to a booked IBNR, and who is waiting on whom

01	**Request the data** Actuarial files a ticket for the quarter's extract	IT
02	**Wait on the extract** The pull sits behind other reporting work	DAYS
03	**Reconcile to the ledger** Paid and incurred tied back to finance, by hand	FINANCE
04	**Rebuild the triangles** Segment mapping and large-loss carve-outs redone each quarter	DAYS
05	**Select factors and methods** Chain ladder, Bornhuetter-Ferguson, Cape Cod, tail	ACTUARIAL
06	**Rebuild the exhibits** Roll-forward and movement memo reformatted by hand	DAYS
07	**Peer review** A second actuary retraces the work from scratch	DAYS
08	**Book and explain** Answer the CFO, the auditor, and the board	DAYS

AFTER · ON TESORA

3-4 hours

Triangles squared on ingest, so the quarter starts at the selection

YOU GIVE The loss run, as it comes out of the system.

YOU GET Reconciled triangles, method comparisons,
and an IBNR with every value traced to source.

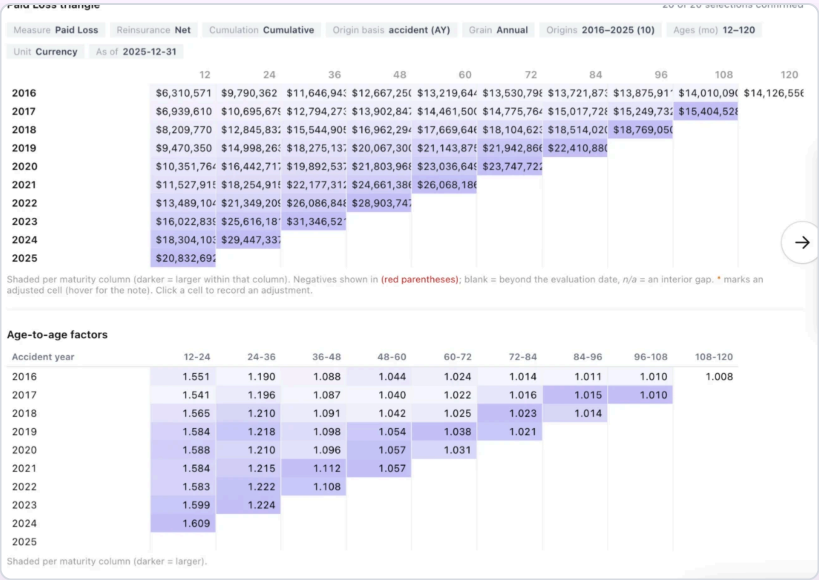

Measure Paid Loss · Reinsurance Net · Cumulation Cumulative · Origin basis accident (AY) · Grain Annual · Origins 2016–2025 (10) · Ages (mo) 12–120 · Unit Currency · As of 2025-12-31

	12	24	36	48	60	72	84	96	108	120
2016	$6,310,571	$9,790,362	$11,646,94	$12,667,25	$13,219,64	$13,530,79	$13,721,87	$13,875,91	$14,010,09	$14,126,55
2017	$6,939,610	$10,695,67	$12,794,27	$13,902,84	$14,461,50	$14,775,76	$15,017,72	$15,249,73	$15,404,52	
2018	$8,209,770	$12,845,83	$15,544,90	$16,962,29	$17,669,64	$18,104,62	$18,514,02	$18,769,05		
2019	$9,470,350	$14,998,26	$18,275,13	$20,067,30	$21,143,87	$21,942,86	$22,410,88			
2020	$10,351,76	$16,442,71	$19,892,53	$21,803,96	$23,036,64	$23,747,72				
2021	$11,527,91	$18,254,91	$22,177,31	$24,661,38	$26,068,18					
2022	$13,489,10	$21,349,20	$26,086,84	$28,903,74						
2023	$16,022,83	$25,616,18	$31,346,52							
2024	$18,304,10	$29,447,33								
2025	$20,832,69									

Shaded per maturity column (darker = larger within that column). Negatives shown in (red parentheses); blank = beyond the evaluation date, *n/a* = an interior gap. * marks an adjusted cell (hover for the note). Click a cell to record an adjustment.

Age-to-age factors

Accident year	12-24	24-36	36-48	48-60	60-72	72-84	84-96	96-108	108-120
2016	1.551	1.190	1.088	1.044	1.024	1.014	1.011	1.010	1.008
2017	1.541	1.196	1.087	1.040	1.022	1.016	1.015	1.010	
2018	1.565	1.210	1.091	1.042	1.025	1.023	1.014		
2019	1.584	1.218	1.098	1.054	1.038	1.021			
2020	1.588	1.210	1.096	1.057	1.031				
2021	1.584	1.215	1.112	1.057					
2022	1.583	1.222	1.108						
2023	1.599	1.224							
2024	1.609								
2025									

Shaded per maturity column (darker = larger).

The full context is baked into the work.

Every session holds the data it used, the assumptions behind it, and the reasoning that got there. Reopen it years later and nothing has to be reconstructed, because the context was never separate from the answer.

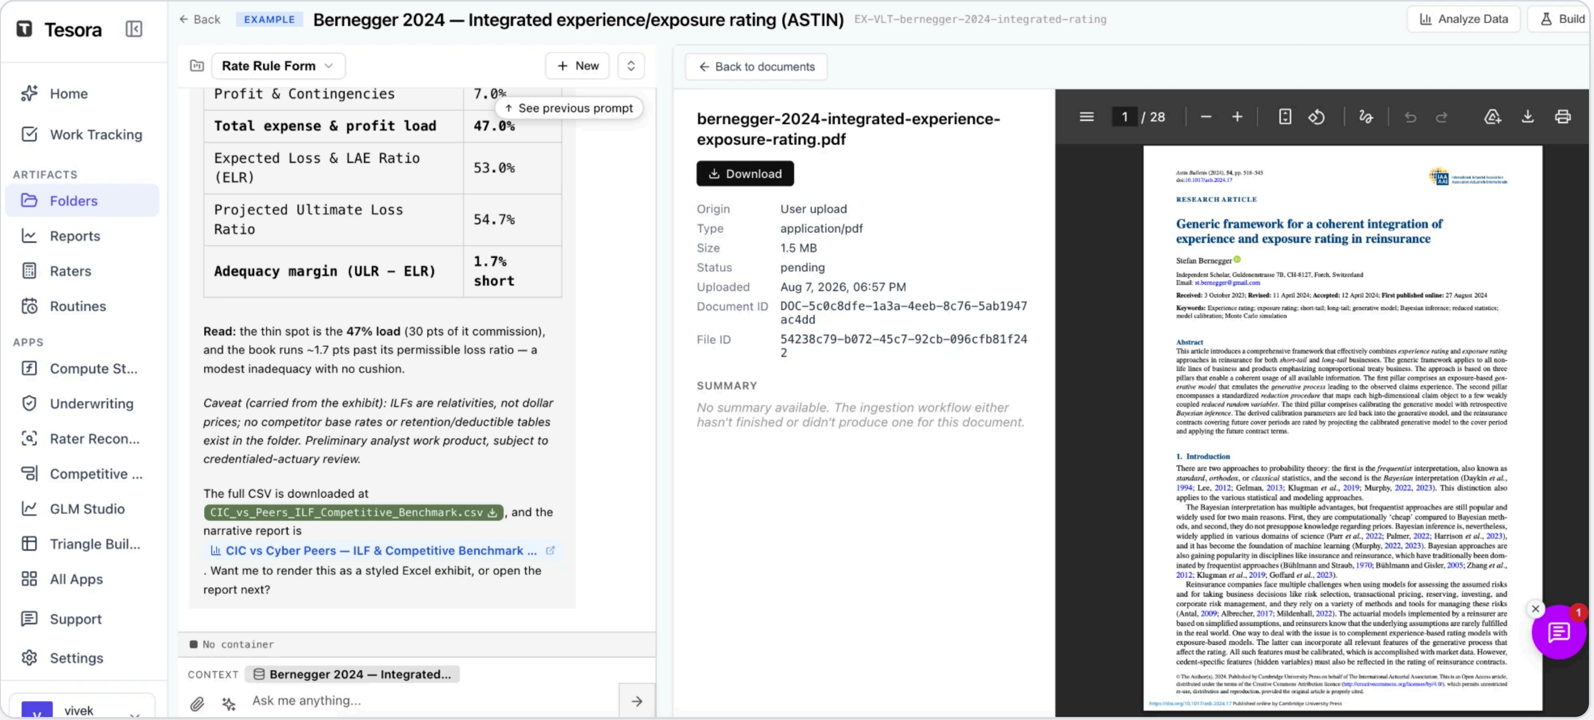

When an actuary leaves

Their assumptions, sources, and reasoning stay in the session. The next person picks up the thread instead of rebuilding it.

Finance

Ties a booked number back to the ledger without waiting on an actuary to explain it.

IT and data

Sees which files fed which number, so a schema change has a known blast radius.

Auditors and regulators

Follow any value to its source document, years after the filing closed.

One shared record of how the work was done, and **it outlasts the people who did it**.

Tesora

Above all, we want to help you win.

We are a thought partner first. This technology is powerful, and it only helps once it is opinionated: built with a real point of view on how the work gets done and defended.

We are here to push the frontier forward with the right partners, not to sell you a license. Doing right by you comes first, so every deployment starts with a proof of concept anchored on your organization's own success metrics. Some teams are stronger at reserving, others at pricing, and we work with yours to find where the workbench sets the best practices and shared learnings first.

Start a POC discussion

sales@tesora.ai

DE-RISKED BY DESIGN

- ✓ It runs in your own environment. We never hold your data.
- ✓ SOC 2 Type II, with the highest data-security controls.
- ✓ We clear your security review before any work touches a filing.

Tesora

APPENDIX

The supporting detail.

We run on top of Claude, and we are model agnostic.

The model is the engine, not the product. What an actuary needs around it is infrastructure: the data wired in, the workflow decided, the cost controlled, and the record kept. That is the part we build, and we build it for you.

01 We guide the model

A raw model will answer anything, confidently. Ours is held to an actuarial method: which factors to consider, which checks to run, when to stop and ask. That guidance is the difference between a plausible answer and a defensible one.

02 We wire the infrastructure

Your loss runs, policy systems, filings, and ledger are connected once and stay connected. A chat window starts empty every time; the workbench already knows your book.

03 The workflow is opinionated

Reserving and pricing have a right order of operations. We encode it, so the work arrives in the shape an actuary expects rather than whatever the prompt happened to produce.

04 We manage the token cost

Running a real book through a frontier model naively is slow and expensive. We route each step to the cheapest model that is still correct, so the economics work at production scale.

We exist to build **your actuarial infrastructure**, not to serve the models or ourselves. When a better model ships, you get it. What stays is the system we built around your book.

It makes us even better.

Better models accelerate the product underneath you, compounding your advantage with every release.

Every release lifts you	The workbench is model-agnostic and runs on the strongest model available. Each new generation handles your schemas and messy edge cases more easily, so the work gets sharper **at no cost to you**.
We keep learning your systems	The longer we work with your team, the more of your data, formats, and quirks the workbench encodes, and that knowledge carries across every model upgrade.
The cheapest correct path	Every org is running a dozen AI pilots. We stay anchored on one thing: the fastest, most correct, and lowest-token way to run a complex workflow within your audit standards.
Audits matter more, not less	As the model does more of the work, proving it matters more, so no step gets skipped and nothing is taken for granted. Every value stays cited, reproducible, and signed off.
The gap widens	Better models unlock even more complex work, so the advantage compounds. A rate change that took months now takes days, so you can reprice against the market while slower competitors wait on a filing cycle. The firms that move first pull ahead, and **standing still gets more expensive**.

The recurring month goes from 150 hours to about 60.

Tesora compresses the mechanical grind, the data pulls, the reconciliations, the model updates, so the same team spends its hours on judgment instead of assembly.

Where the recurring month goes — reclaimed / still hands-on — ILLUSTRATIVE · MODELED

Data extraction & reconciliation	39 → 4	10×
Pricing model updates	28 → 7	4×
Loss trend analysis	23 → 12	2×
Reserve & loss development	21 → 12	1.7×
Management & board reporting	17 → 10	1.7×
Reinsurance program analysis	12 → 4	3×
Peer review & documentation	10 → 8	1.3×
Monthly total	150 → ~60	2.5×

2.5×

the recurring workload, same team and sign-off

The blended throughput gain across the whole recurring month.

Smarter not harder

Time freed for strategy and selections, not assembly

Unlock your actuaries

More products, territories, and actuarial capability

Hours reflect a typical ~150 hour recurring actuarial month, with per-workflow gains from Tesora's engagement framework. Figures are modeled and confirmed against measured results during a build.

We hand the mechanical hours back to your team as actuarial judgment, freeing them to focus where it counts. Capacity stops being the ceiling.

A team built at the intersection of frontier AI research and decades of actuarial experience, the two disciplines this problem demands.

Vivek Rao

CEO

Former private-equity investor and McKinsey consultant focused on P&C insurance.

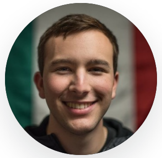

Federico Reyes

CTO

Stanford AI researcher. Graph neural nets and relational data at Kumo AI (acquired by Nvidia), and Google's Document OCR team.

Philo Bishay

CPO & CHIEF ACTUARY

Chief actuary at W. R. Berkley E&S, owning pricing and reserving. Capital modeler at **Gen Re**.

MP

Michael Palmich

HEAD OF SALES

Built and sold actuarial software companies for two decades, including Arius and Akur8.

EL

Eric Lam

FOUNDING ENGINEER

MIT math, CS, and physics, with **three actuarial exams**. Writes the kernels.

ADVISORY BOARD

Ethan Genteman
Former Chief Actuary, **Palomar**

Joanne Chen
Partner, **Foundation Capital**

Actuaries who have signed the filings, and investors who have built the category, keep us honest about what carriers will actually adopt.

BACKED BY

foundation capital — Leading AI and enterprise investors since 1995

Y Combinator S25
The accelerator behind Stripe and Airbnb

☆ Proud sponsors of CAS