



CASE STUDY · SPECIALTY P&C

Frontier AI for Actuaries.

One specialty carrier moved pricing, reserving, and rate filing onto the Tesora Workbench, with every number traced back to the document it came from.

ENGAGEMENT

End-to-end rebuild

SEGMENT

Admitted specialty P&C

PLATFORM

Tesora Workbench

THE ENGAGEMENT

A US admitted specialty P&C carrier

A full rebuild of pricing, reserving, and rate filing on the Tesora Workbench. Carrier named on request, under NDA.

SCOPE

Pricing, reserving, filing

DATE

July 2026

BACKED BY

Foundation Capital · Y Combinator

CONFIDENTIAL · PROVIDED UNDER NDA



The Workbench is a chat surface that sits next to Excel. You brief it the way you would a senior analyst, four agents do the work and cite every number, and the answer comes back as a spreadsheet or as runnable code. Five capabilities ran on it in this engagement.

Tesora Workbench

ASK

Indicate GA commercial auto, on-leveled, with exhibits, and update the rater.

Run

WORKING, EVERY STEP TRACED

✓

Ingestion Agent read the GA loss runs, 2019 to 2025
↳ your loss data, 41,208 records

✓

Insight Agent fit severity and frequency, trend 6.2%
↳ cell C42

✓

Rating Agent on-leveled to current rates, factor 1.18
↳ SERFF GA-2024-00317

✓

Audit Agent logged every value and flagged two capped claims
↳ append-only audit trail

Excel

indication and exhibits

Code

the rater, versioned

THE FIVE CAPABILITIES

01 Competitive Market Analysis

Any admitted SERFF filing benchmarked against your rates, by class and territory.

02 Rate Impact Studies

Indicated versus selected and premium impact across the whole book before you file.

03 Loss Modeling for Pricing

Frequency, severity, trend, and on-leveling, assembled so the week goes to selections.

04 Reserving Analysis

Loss development with the triangles pulled and reconciled before you sit down.

05 Rater Deployment

Tested raters live as an API or GUI, with filing support for all 50 states.



The Workbench is the surface. The Tesora **Harness** is the framework beneath it that turns a fast answer into one you can put your name on. Four agents do the work, and not one of them gets the last word.

Ingestion Agent **READS**

Takes loss runs, SOVs, dec pages, and rate manuals in whatever format they arrive, and normalizes them underneath the workflow you already run.

Insight Agent **MODELS**

Fits the curves, trends, and on-levels, then hands back a write-up the reviewer can check line by line.

Rating Agent **DEPLOYS**

Turns the actuary's selections into a tested rater and the state exhibits that go with it, ready to file.

Audit Agent **PROVES**

Traces every value to a page, a row, a cell, and writes down anything it skipped, so a capped limit lands in the memo where a reviewer will see it.

Author, reviewer, and approver stay three separate people, enforced. Every value cites its source, the log is append-only, and **the actuary owns every decision.**

SOC 2 ready

SOX controls

ASOP 23 / 41 / 56

Single-tenant VPC

Zero data retention

	Actuarial Consultants Milliman · WTW · Oliver Wyman	Your Team, By Hand Credentialed actuaries, manual	DIY / General AI ChatGPT · Claude Code · copilots	T Tesora Trained actuarial system
Turnaround	✗ Weeks to months; you wait in their queue	✗ Bounded by headcount; the backlog never clears	≈ Instant, but the answer is not usable	✓ Minutes to days, inside your workflow
Cost model	✗ \$300 to \$800 per hour; six-figure projects	≈ Fixed seats; every project competes for them	≈ Cheap per query; rework erases the savings	✓ Per-seat plus usage; scales with the book
Accuracy & sourcing	≈ High quality, delivered as a static report	≈ Trusted, but slow and manual to reproduce	✗ Hallucinated numbers, no source citations	✓ Cell-level citations to your data, every value
Fits your format & filings	✗ Generic deliverables, not your rater	≈ Your format, rebuilt by hand each time	✗ No rater, no exhibits, nothing fileable	✓ Your raters and exhibits, filing-ready in 50 states
Audit trail	≈ Their workpapers, not yours	≈ Manual side-by-side QA against source	✗ None; a number cannot be traced back	✓ Append-only log; author, reviewer, approver enforced

✓

Strength

≈

Partial or conditional

✗

Gap



WHERE THE ACTUARIAL HOURS WENT

✗ **Raters versioned by filename**

The rating logic lived in spreadsheets versioned by filename, which is to say not versioned at all. Changing one meant an IT ticket and a wait.

✗ **One team, every priority at once**

Pricing, competitive work, reserving, and rate filing all drew on the same handful of credentialed actuaries. Whatever was due soonest won.

✗ **Filings built by hand, a state at a time**

Each rate filing meant weeks of assembling exhibits and reformatting them to the next state's template.

✗ **No quick read on the market**

Comparing the carrier's rates to competitor SERFF filings happened rarely, and when it did it took days.

THE REAL COST

None of this was a headcount problem, it was a **throughput** problem. The hours went to reconciling data and reformatting exhibits. The selections that actually move a loss ratio came last, when they came at all.

The carrier was **capacity-bound, not market-bound**. The demand for pricing work was there. The hours to do it were not.

~150 hrs_{/mo}

Recurring actuarial workflow before Tesora, most of it spent moving and reconciling data before any analysis began.



The rebuild ran in three moves, all inside the carrier's own environment, with a founder on the other end of the phone the whole way.

1 INGEST

Mass data ingestion

Every policy, loss run, and rate manual loaded and normalized under the existing workflow, in whatever format it arrived.

- Any carrier format, read as it is
- The full in-force book and its loss history
- Anomalies flagged before they reached a model

2 REBUILD AND RUN

Every rater run across the book

The Ingestion and Rating Agents rebuilt the raters as versioned, tested models and repriced the entire in-force book.

- Existing rating logic preserved exactly
- Every in-force policy repriced, the full book
- Version history and test cases from day one

3 OPERATE

The five capabilities, live

Pricing, reserving, competitive analysis, and filing now run on one surface. They used to compete for the same few desks.

- One append-only audit trail
- The actuary signs off on every output
- New analyses in minutes, not weeks

- ✓ Single-tenant VPC
- ✓ No data leaves the boundary
- ✓ Author, reviewer, approver enforced
- ✓ Append-only audit trail from day one


60-day build → production

WHAT CHANGED, CAPABILITY BY CAPABILITY

Competitive Market Analysis Before, competitor filings were read by hand, days per comparison. Now, any admitted filing is benchmarked in minutes, class by class.
Rate Impact Studies Before, impact was estimated late, on a subset. Now, indicated versus selected runs across the whole book, every driver traced.
Loss Modeling for Pricing Before, curves and on-leveling ate the analyst's week. Now, the fits are assembled and the week goes to the selections.
Reserving Analysis Before, reserving started with days of reconciling triangles. Now, the triangles are squared before the analysis begins.
Rater Deployment & Filing Before, a change waited on IT and each state was built by hand. Now, it goes live as an API, with exhibits for all 50 states.

PROOF · ONE LINE, ONE LOOK

The competitive read the carrier had never had time to run: where its loss cost sits against the market it files in.



ILLUSTRATIVE The carrier sat **6% below market** on this line, and had never had a fast way to see it.

WHERE THE HOURS WENT, AND THE MULTIPLIER THE WORKBENCH DELIVERED ON EACH

ILLUSTRATIVE · MODELED

RECURRING ACTUARIAL WORKFLOW	HRS/MO	TESORA
Data extraction & reconciliation	39	10×
Pricing model updates	28	4×
Loss trend analysis	23	2×
Reserve & loss development	21	1.7×
Management & board reporting	17	1.7×
Reinsurance program analysis	12	3×
Peer review & documentation	10	1.3×
Monthly total	150	2.5× avg

2.5×

More recurring work, same team

Average throughput gain across the monthly actuarial workflow.

4.4×

Faster rate-filing prep

Loss costs, exhibits, and state formatting, on the project work.

5×

Competitor analysis

Ingesting and structuring competitor filings, minutes not days.

30

Books one actuary can cover

Up from roughly five, at the same headcount and sign-off.

Multipliers reflect Tesora's engagement framework applied to the carrier's stated ~150 hour monthly baseline. Final figures to be confirmed against measured results before external use.

Same team and the same sign-off, with close to an order of magnitude more work through the door. The carrier stopped being capacity-bound.



The carrier rebuilt its actuarial function on Tesora and kept **every number defensible.**

It started with mass data ingestion and every rater run across the book, then pricing, reserving, competitive work, and 50-state filing moved onto the Workbench. Consultants had always traded speed for cost, the in-house team traded reach for hours, and general AI traded trust for speed. **Tesora gave the carrier all three at once**, which is the hard part, and the team took on more without adding a seat.

RECURRING THROUGHPUT	2.5×
RATE-FILING PREP	4.4×
FILING COVERAGE	50 states
AUDIT TRAIL	cell-level

Talk to Tesora.

60-day build in your own environment. Sign or refund.

founders@tesora.ai

tesora.ai · Frontier AI for Actuaries