



CONFIDENTIAL

FRONTIER AI FOR ACTUARIES

One actuary, the reach of a whole team.

Tesora rebuilds specialty P&C pricing, reserving, and rate filing on AI that reads your data, shows its work, and lets you file with your name on it.

FOCUS

US specialty P&C

PRODUCT

The Tesora Workbench

BACKED BY

Foundation Capital · Y Combinator

WHO IS BUILDING THIS

The big labs don't understand **actuarial work**.

A general model is built to be broadly capable. A rate filing needs the opposite: every number traced to its source, nothing invented, and a paper trail an examiner will accept. Tesora is AI researchers from Stanford, Google, and Kumo who went deep on that problem, next to a chief actuary who has signed the filings. Backed by **Foundation Capital** and **Y Combinator**.



Vivek Rao

CO-FOUNDER & CEO

McKinsey, on insurance and PE. Then a PE investor at Sycamore, behind the Walgreens buyout. Runs the customer engagements and go-to-market.



Federico Reyes Gómez

CO-FOUNDER & CTO

Stanford AI researcher. Built Document AI at Google and predictive models at Kumo, acquired by Nvidia. Leads the research.



Philo Bishay

CO-FOUNDER & CPO

Former chief actuary. Priced and reserved specialty P&C at GenRe and Berkley Specialty.



WHERE THE MONTH GOES

Most of the actuarial month is spent **getting data ready**.

A mid-sized specialty team runs roughly 150 hours a month of recurring actuarial work. The majority of it goes to reconciling data and reformatting exhibits. The selections that actually move a loss ratio come last.

Every hour spent on intake is an hour taken from the market, the mix, and the next filing.

RECURRING ACTUARIAL MONTH · HOURS BY TASK



The single biggest block is **data extraction and reconciliation**, the least actuarial part of an actuary's month.



Three costs follow from that, and every specialty carrier feels all three.

RATE FILING

Filings lag the market

Each state filing is weeks of assembling exhibits and reformatting them to the next template. By the time a rate is live, the trend it was built on has moved.

Rate adequacy erodes before it goes live.

COMPETITIVE READ

The market is a blind spot

Competitor SERFF filings rarely get read, because reading them by hand takes days. So carriers price against their own history instead of the market they compete in.

You price on your own book alone.

CAPACITY

The team is the ceiling

Every new line, state, or study competes for the same few credentialed actuaries. Growth means hiring, and credentialed hires are slow and scarce.

Growth is capped at your hiring rate.



The three ways carriers try to fix this each hit a wall.

ACTUARIAL CONSULTANTS

- ✗ Weeks to months per project, at **\$300 to \$800 an hour**, and you wait in their queue.
- ✗ You get a static report. No rater you can run, no state templates you can file.
- ✗ The model does not scale. You buy hours, and hours run out.

THE TEAM, BY HAND

- ✗ Trusted, but **bounded by headcount**, and the backlog never fully clears.
- ✗ Every change is rebuilt and re-checked by hand, which is slow to reproduce.
- ✗ The most senior people spend their week on the least senior work.

DIY / GENERAL AI

- ✗ Fast and cheap, but it **hallucinates numbers** with no source you can cite.
- ✗ No rater, no exhibits, nothing an actuary can actually file.
- ✗ Ask twice, get two answers, which is disqualifying in a rate filing.

Consultants and your team run out of hours, and general AI never had the trust to begin with. That gap is the whole opportunity.

The Workbench is a chat surface that sits next to Excel. You brief it the way you would a senior analyst, four agents do the work and cite every number, and the answer comes back as a spreadsheet or as runnable code.

T

Tesora Workbench

ASK

Indicate GA commercial auto, on-leveled, with exhibits, and update the rater.

Run

WORKING, EVERY STEP TRACED

✓

Ingestion Agent

read the GA loss runs, 2019 to 2025

↳ your loss data, 41,208 records

✓

Insight Agent

fit severity and frequency, trend 6.2%

↳ cell C42

✓

Rating Agent

on-leveled to current rates, factor 1.18

↳ SERFF GA-2024-00317

✓

Audit Agent

logged every value and flagged two capped claims

↳ append-only audit trail

Excel

indication and exhibits

Code

the rater, versioned

WHAT RUNS ON IT

- 01

Competitive Market Analysis

Any admitted SERFF filing benchmarked against your rates, by class and territory.
- 02

Rate Impact Studies

Indicated versus selected and premium impact across the whole book before you file.
- 03

Loss Modeling & Reserving

Fits, trend, on-leveling, and loss development, with the data prep already done.
- 04

Rater Deployment & Filing

Tested raters live as an API or GUI, with filing support for all 50 states.



The Workbench is the surface. The Tesora **Harness** is the framework beneath it that turns a fast answer into one you can put your name on. Four agents do the work, and not one of them gets the last word.

Ingestion Agent **READS**

Takes loss runs, SOVs, dec pages, and rate manuals in whatever format they arrive, and normalizes them underneath the workflow you already run.

Insight Agent **MODELS**

Fits the curves, trends, and on-levels, then hands back a write-up the reviewer can check line by line.

Rating Agent **DEPLOYS**

Turns the actuary's selections into a tested rater and the state exhibits that go with it, ready to file.

Audit Agent **PROVES**

Traces every value to a page, a row, a cell, and writes down anything it skipped, so a capped limit lands in the memo where a reviewer will see it.

Author, reviewer, and approver stay three separate people, enforced. Every value cites its source, the log is append-only, and the actuary owns every decision.

SOC 2 ready

SOX controls

ASOP 23 / 41 / 56

Single-tenant VPC

Zero data retention

One carrier now does the work of five, at the same headcount.

A US admitted specialty carrier rebuilt pricing, reserving, and rate filing on the Workbench. The team did not grow. What it could get through the door did.

2.5×

More recurring work, same team

4.4×

Faster rate-filing prep

50

States with filing support

100%

Values traced to source

THE CHALLENGE

Capacity-bound, not market-bound

Raters lived in spreadsheets versioned by filename. Pricing, reserving, competitive work, and filing all drew on the same few actuaries, and whatever was due soonest won.

WITH TESORA

Five kinds of work, in parallel

After mass data ingestion and every rater run across the book, the five workstreams now run at once on one platform, and **every number still traces to a source**. Built inside the carrier's own environment, on their own data.



FRONTIER AI FOR ACTUARIES

Built by the people who **priced these lines.**

Pricing, reserving, and rate filing, fast enough to keep up with the market and defensible enough to file.

Talk to Tesora.

We build it on your own data, inside your own environment.

founders@tesora.ai

tesora.ai · US specialty P&C